

TP4 - OpenMP (Parallel Sections, Single, Master, Synchronization)

Imad Kissami

February, 2025

Exercise 1 : Work Distribution with Parallel Sections

- Write a program that initializes an array of size N with random values.
- Use `#pragma omp sections` to divide the work :
 - Section 1 : Compute the sum of all elements.
 - Section 2 : Compute the maximum value.
 - Section 3 : Compute the standard deviation.
- Ensure that all computations run in parallel.
- Do not recompute the sum inside Section 3. The standard deviation must use the result obtained in Section 1.

```

1 #include <stdio.h>
2 #include <stdlib.h>
3 #include <math.h>
4 #include <time.h>
5
6 #define N 1000000
7
8 int main() {
9
10    double *A = malloc(N * sizeof(double));
11    if (A == NULL) {
12        printf("Memory allocation failed\n");
13        return 1;
14    }
15
16    double sum = 0.0;
17    double mean = 0.0;
18    double stddev = 0.0;
19    double max;
20
21    // Initialization
22    srand(0);
23    for (int i = 0; i < N; i++)
24        A[i] = (double)rand() / RAND_MAX;
25
26    // Compute sum and max
27    sum = 0.0;
28    max = A[0];
29    for (int i = 0; i < N; i++) {
30        sum += A[i];
31        if (A[i] > max)
32            max = A[i];

```

```

33     }
34
35     // Compute mean
36     mean = sum / N;
37
38     // Compute standard deviation
39     stddev = 0.0;
40     for (int i = 0; i < N; i++)
41         stddev += (A[i] - mean) * (A[i] - mean);
42
43     stddev = sqrt(stddev / N);
44
45     // Print results
46     printf("Sum=%f\n", sum);
47     printf("Max=%f\n", max);
48     printf("Std Dev=%f\n", stddev);
49
50     free(A);
51     return 0;
52 }
```

Exercise 2 : Exclusive Execution - Master vs Single

- Write a program where :
- A master thread initializes a matrix.
- A single thread prints the matrix.
- All threads compute the sum of all elements in parallel.
- Compare execution time with and without OpenMP.

```

1 #include <stdio.h>
2 #include <stdlib.h>
3 #include <time.h>
4
5 #define N 1000
6
7 void init_matrix(int n, double *A) {
8     for (int i = 0; i < n; i++) {
9         for (int j = 0; j < n; j++) {
10             A[i*n + j] = (double)(i + j);
11         }
12     }
13 }
14
15 void print_matrix(int n, double *A) {
16     for (int i = 0; i < n; i++) {
17         for (int j = 0; j < n; j++) {
18             printf("%6.1f ", A[i*n + j]);
19         }
20         printf("\n");
21     }
22 }
23
24 double sum_matrix(int n, double *A) {
25     double sum = 0.0;
26     for (int i = 0; i < n*n; i++) {
27         sum += A[i];
28     }
29     return sum;
30 }
31
32 int main() {
33
34     double *A;
35     double sum;
36     double start, end;
37
38     A = (double*) malloc(N * N * sizeof(double));
```

```

39     if (A == NULL) {
40         printf("Memory allocation failed\n");
41         return 1;
42     }
43
44     start = (double) clock() / CLOCKS_PER_SEC;
45
46     /* Initialization */
47     init_matrix(N, A);
48
49     /* Printing (can be commented if N is large) */
50     /* print_matrix(N, A); */
51
52     /* Sum computation */
53     sum = sum_matrix(N, A);
54
55     end = (double) clock() / CLOCKS_PER_SEC;
56
57     printf("Sum=%lf\n", sum);
58     printf("Execution time=%f seconds\n", end - start);
59
60     free(A);
61     return 0;
62 }
```

Exercise 3 : Load Balancing with Parallel Sections

- Implement a task scheduling mechanism using parallel sections.
- Simulate three different workloads :
 - Task A (light computation)
 - Task B (moderate computation)
 - Task C (heavy computation)
- Measure the execution time and optimize the workload distribution.

```

1 #include <math.h>
2
3 void task_light(int N) {
4     double x = 0.0;
5     for (int i = 0; i < N; i++) {
6         x += sin(i * 0.001);
7     }
8 }
9
10 void task_moderate(int N) {
11     double x = 0.0;
12     for (int i = 0; i < 5*N; i++) {
13         x += sqrt(i * 0.5) * cos(i * 0.001);
14     }
15 }
16
17 void task_heavy(int N) {
18     double x = 0.0;
19     for (int i = 0; i < 20*N; i++) {
20         x += sqrt(i * 0.5) * cos(i * 0.001) * sin(i * 0.0001);
21     }
22 }
```

Exercise 4 : Synchronization and Barrier Cost

- Implement a dense matrix-vector multiplication.
- Version 1 : Implicit barrier.
- Version 2 : Use `schedule(dynamic)` with `nowait`.

- Version 3 : Use `schedule(static)` with `nowait`.
- Measure :
 - CPU time
 - Speedup
 - Efficiency
- Explain :
 - Why barriers limit scalability.
 - When `nowait` becomes dangerous.
- Run the code with 1, 2, 4, 8, 16 threads and compare the execution time of version 1 and 3. Plot for each version the :
 - CPU time
 - Speedup & Efficiency
 - $\text{MFLOP/s} = \frac{\text{FLOPs}}{t \times 10^6}$

```

1 #include <stdio.h>
2 #include <stdlib.h>
3 #include <omp.h>
4
5 void dmvm(int n, int m,
6 double *lhs, double *rhs, double *mat)
7 {
8     for (int c = 0; c < n; ++c) {
9         int offset = m * c;
10        for (int r = 0; r < m; ++r)
11            lhs[r] += mat[r + offset] * rhs[c];
12    }
13 }
14
15 int main(void)
16 {
17     const int n = 40000; // columns
18     const int m = 600; // rows
19
20     double *mat = malloc(n * m * sizeof(double));
21     double *rhs = malloc(n * sizeof(double));
22     double *lhs = malloc(m * sizeof(double));
23
24     // initialization
25     for (int c = 0; c < n; ++c) {
26         rhs[c] = 1.0;
27         for (int r = 0; r < m; ++r)
28             mat[r + c*m] = 1.0;
29     }
30
31     for (int r = 0; r < m; ++r)
32         lhs[r] = 0.0;
33
34     dmvm(n, m, lhs, rhs, mat);
35
36     // print result
37     printf("lhs:\n");
38     for (int r = 0; r < m; ++r)
39         printf("lhs[%d] = %.1f\n", r, lhs[r]);
40
41
42     free(mat);
43     free(rhs);
44     free(lhs);
45     return 0;
46 }
```